

Fully connected neural network on MNIST dataset

- a) Open the notebook [fcn_MNIST.ipynb](#). In this notebook we use a fully connected neural network to predict the handwritten digits of the MNIST dataset.
We have 4000 examples with 784 pixel values and 10 classes. Run the first 3 cells.
- b) Write the missing TensorFlow code in cell 4 for the first hidden layer.
- c) Run the next two cells to store the graph and do a forward pass of the untrained network. Have a look at the network.
- d) Now let's train the model. Finish the code to calculate the loss and accuracy for the validation set.

Fully connected neural network on MNIST dataset (Tricks)

Note for docker users.

- In this notebook we create different runs so it might be beneficial to save them also outside the docker container. This is possible using the `-v` option when starting docker.

```
docker run -p 8888:8888 -p 6006:6006 -v /Users/oli/Documents/workspace/dl_course/:/notebooks
```

- If you experience crashes of the docker container do a two step procedure.
First start docker in bash.

```
docker run -p 8888:8888 -p 6006:6006 -v /Users/oli/Documents/workspace/dl_course/:/notebooks
```

Then start the jupyter notebook in the console with

```
jupyter notebook --NotebookApp.token=tensorchiefs
```

- a) Open the notebook [fcn_MNIST_keras](#) and run the first model (execute the cell after training) and visualize the result in TensorBoard (have a look at learning curves and the histograms / distributions of the weights)
- b) Remove the `init='zero'` argument of the dense layers, to have a proper initialization of your weights. Change the name from `name = 'sigmoid_init0'` to `name = 'sigmoid'`. Restart the kernel and repeat the training as in a). Compare the results in TensorBoard, describe your results.

- c) Change the activations / non-linearities from `Activation('sigmoid')` to `Activation('relu')` and change the name from `name = 'sigmoid'` to `name = 'relu'`. Continue as above, especially have a look at the validation loss do you observe overfitting.
- d) Add a dropout layer: Now add a dropout layer `model.add(Dropout(0.3))` between the Dense-Layer and the Activation. Change the name from `name = 'relu'` to `name = 'dropout'`.
- e) Add a batch-normalization: Now add a batch-norm layer `model.add(BatchNormalization())` between the Dense-Layer and the Dropout. Change the name from `name = 'dropout'` to `name = 'batch_dropout'`. Continue as above

The network should look like:

Layer (type)	Output Shape	Param #	Connected to
dense_1 (Dense)	(None, 500)	392500	dense_input_1[0][0]
batchnormalization_1 (BatchNormaliza	(None, 500)	2000	dense_1[0][0]
dropout_1 (Dropout)	(None, 500)	0	batchnormalization_1[0][0]
activation_1 (Activation)	(None, 500)	0	dropout_1[0][0]
dense_2 (Dense)	(None, 50)	25050	activation_1[0][0]
batchnormalization_2 (BatchNormaliza	(None, 50)	200	dense_2[0][0]
dropout_2 (Dropout)	(None, 50)	0	batchnormalization_2[0][0]
activation_2 (Activation)	(None, 50)	0	dropout_2[0][0]
dense_3 (Dense)	(None, 10)	510	activation_2[0][0]
Total params: 420,260			
Trainable params: 419,160			
Non-trainable params: 1,100			