

Deep Learning in MSE MachLe

Convolutional Neural Networks for image data

Exercise 1:

We want to investigate, if a CNN outperforms a fc NN on image data.

First we recall the design of the fc NN which performed so far best on MNIST when only keeping 4000 examples in the training data set (see below). With this NN we have reached ~91% accuracy on the validation data set.

Layer (type)	Output Shape	Param #
dense_1 (Dense)	(None, 500)	392500
batch_normalization_1 (Batch Normalization)	(None, 500)	2000
dropout_1 (Dropout)	(None, 500)	0
activation_1 (Activation)	(None, 500)	0
dense_2 (Dense)	(None, 50)	25050
batch_normalization_2 (Batch Normalization)	(None, 50)	200
dropout_2 (Dropout)	(None, 50)	0
activation_2 (Activation)	(None, 50)	0
dense_3 (Dense)	(None, 10)	510
Total params: 420,260.0		
Trainable params: 419,160.0		
Non-trainable params: 1,100.0		

- a) Where do we spend most learnable parameter? Can you explain the "Param #" of the dense_1 layer?

Remark: dense layer is the same as fully connected layer.

- b) Now we want to use our first CNN with only 1 convolutional and 1 dense layer:

Layer (type)	Output Shape	Param #	Connected to
convolution2d_1 (Convolution2D)	(None, 28, 28, 32)	320	convolution2d_input_1[0][0]

activation_1 (Activation)	(None, 28, 28, 32)	0	convolution2d_1[0][0]
flatten_1 (Flatten)	(None, 25088)	0	activation_1[0][0]
dense_1 (Dense)	(None, 10)	250890	flatten_1[0][0]
activation_2 (Activation)	(None, 10)	0	dense_1[0][0]
=====			
Total params: 251,210			
Trainable params: 251,210			
Non-trainable params: 0			

In which layer do we need to learn most parameter/weights?

Do you expect with this cnn1 more or less overfitting then in the fc NN above? Why?

Please open the ipython-Notebook 08_cnn1_mnist.ipynb and try to understand the code and run the code.

How large is the accuracy on the validation set? Do you observe overfitting? Describe how you check for overfitting and/or sketch the corresponding graph.

Restart the kernel and run the model without first standardizing the data which is done in code cell 7 after # here we center and standardize the data. Instead of commenting out each command, you can turn the cell in Cell Type "markdown" and then run the code.

How large is now the accuracy on the validation set? Can you explain, what happened?

- c) Please open the ipython-Notebook 08_cnn2_mnist_No_solution.ipynb. Search for the position in the code which is marked by "# here is your code coming:" and add the missing layers - the missing layers are marked below. How is the performance of cnn2?

Layer (type)	Output Shape	Param #	Connected to
=====			
convolution2d_1 (Convolution2D)	(None, 28, 28, 32)	320	convolution2d_input_1[0][0]
batchnormalization_1 (BatchNorma	(None, 28, 28, 32)	128	convolution2d_1[0][0]
activation_1 (Activation)	(None, 28, 28, 32)	0	batchnormalization_1[0][0]
convolution2d_2 (Convolution2D)	(None, 28, 28, 32)	9248	activation_1[0][0]
batchnormalization_2 (BatchNorma	(None, 28, 28, 32)	128	convolution2d_2[0][0]
activation_2 (Activation)	(None, 28, 28, 32)	0	batchnormalization_2[0][0]
maxpooling2d_1 (MaxPooling2D)	(None, 14, 14, 32)	0	activation_2[0][0]
convolution2d_3 (Convolution2D)	(None, 14, 14, 64)	18496	maxpooling2d_1[0][0]
batchnormalization_3 (BatchNorma	(None, 14, 14, 64)	256	convolution2d_3[0][0]
activation_3 (Activation)	(None, 14, 14, 64)	0	batchnormalization_3[0][0]
-----START OF THE MISSING LAYERS-----			
convolution2d_4 (Convolution2D)	(None, 14, 14, 64)	36928	activation_3[0][0]

batchnormalization_4 (BatchNorma	(None, 14, 14, 64)	256	convolution2d_4[0][0]
activation_4 (Activation)	(None, 14, 14, 64)	0	batchnormalization_4[0][0]
maxpooling2d_2 (MaxPooling2D)	(None, 7, 7, 64)	0	activation_4[0][0]
-----END OF THE MISSING LAYERS-----			
flatten_1 (Flatten)	(None, 3136)	0	maxpooling2d_2[0][0]
dense_1 (Dense)	(None, 200)	627400	flatten_1[0][0]
batchnormalization_5 (BatchNorma	(None, 200)	800	dense_1[0][0]
dropout_1 (Dropout)	(None, 200)	0	batchnormalization_5[0][0]
activation_5 (Activation)	(None, 200)	0	dropout_1[0][0]
dense_2 (Dense)	(None, 10)	2010	activation_5[0][0]
activation_6 (Activation)	(None, 10)	0	dense_2[0][0]
=====			
Total params: 695,970			
Trainable params: 695,186			
Non-trainable params: 784			

8 faces fine tuning

In this exercise we work with the 8 faces dataset. We want to improve the performance by using a pretrained vgg16 network. We predict the features on the fc1 layer with the already learned weights on imagenet and then train a small fully connected network for our own labels. The feature extraction was done in this notebook [vgg16_feature_extraction_8_faces](#)

- What do you expect, will it increase our performance? Why? What's the idea behind this so called fine tuning?
- Open the notebook [8 faces fine tuning](#) and build this network and then train it.

Layer (type)	Output Shape	Param #	Connected to
=====			
dense_1 (Dense)	(None, 400)	1638800	dense_input_1[0][0]
batchnormalization_1 (BatchNorma	(None, 400)	1600	dense_1[0][0]
activation_1 (Activation)	(None, 400)	0	batchnormalization_1[0][0]
dropout_1 (Dropout)	(None, 400)	0	activation_1[0][0]
dense_2 (Dense)	(None, 400)	160400	dropout_1[0][0]
batchnormalization_2 (BatchNorma	(None, 400)	1600	dense_2[0][0]
activation_2 (Activation)	(None, 400)	0	batchnormalization_2[0][0]
dropout_2 (Dropout)	(None, 400)	0	activation_2[0][0]
dense_3 (Dense)	(None, 8)	3208	dropout_2[0][0]
activation_3 (Activation)	(None, 8)	0	dense_3[0][0]

```
=====
Total params: 1,805,608
Trainable params: 1,804,008
Non-trainable params: 1,600
=====
```

- c) Complete the code to get the predicted labels out of the probability vector and look at the accuracy on the test data.